Rules for using AI

An adult guide to safer decisions

14 rules for personal life, work and the information you share

14

rules for everyday use

Adult Edition

October 2026 / Layout edition 1.1

Find your rule

Select any title to go straight to the page.

How to use this book 4

Personal

1. Share the question before the statement 5

Finance and privacy

2. Never buy certainty from a prediction 6

Finance and investing

3. Verify the person before moving the money 7

Money and identity

4. Let AI prepare questions for care 8

Health

5. Do not outsource your support network 9

Wellbeing and relationships

6. Read the policy behind the promise 10

Shopping and travel

Find your rule

Select any title to go straight to the page.

Work

7. Permission comes before the upload 11

Confidential information

8. Open every source that carries the argument 12

Research and legal information

9. Keep a person accountable for every draft 13

Writing and applications

10. Give automation only the access it needs 14

Tools and automation

Public life

11. Check the claim before you amplify it 15

News and social media

12. Ask before borrowing a face or voice 16

Images and consent

Family

13. Agree on boundaries before there is a problem 17

Supporting teenagers

Everyday practice

14. Measure the result before trusting the shortcut 18

Choosing useful tools

Your check before sending 19

Conclusion 20

How to use this book

AI can help you plan, learn, draft and compare. It can also give you a convincing wrong answer or invite you to share more information than a task needs. The aim of this book is to help you use it deliberately and keep responsibility for the result.

The structure

Choose a category in the contents, then a subcategory and a numbered rule. Each rule gives you a do, a don’t, evidence or context, a practical suggestion and a conclusion. These are memorable working principles, not legal rules.

Read the evidence carefully

News sources are marked N; official guidance is marked G. Sources and further reading appear beneath each rule. Click a reference or article title to open the original. An incident illustrates a risk; it does not prove that every AI product behaves the same way. Recommendations and fictional examples are identified separately from reported facts.

Know the limits

This guide concerns everyday generative AI and chatbots. Specialized, regulated services may have different protections. Before using sensitive data, check the exact service, account, contract and permissions. Financial, medical and legal sections support better questions and verification; individual decisions may require a qualified professional.

Navigation

Select a title in the contents to jump to it. Each chapter has a return link. The Word edition uses stable chapter links; page numbers can change when you edit. Sources were checked on 7 October 2026. Historical incidents are dated, not presented as current product tests.

01 Share the question before the statement

Personal / Finance and privacy

A budget question rarely needs your complete financial history. A statement can reveal account numbers, addresses, pay dates and transactions involving other people. Begin by deciding what information is truly necessary.

Do

Use invented figures or a short category summary. For a real calculation, keep identifying information out and check the arithmetic yourself. Use your bank's own secure tools when the task requires access to an account.

Don't

Upload full bank statements, payment-card images, tax identifiers, passwords or one-time codes to a general-purpose chatbot. Covering your name alone does not remove the rest of the identifying detail.

News and its limits

Forbes reported in May 2023 that Samsung restricted staff use of generative AI after sensitive code was uploaded. [N1] This was a workplace incident, not a reported bank-statement theft. The broader lesson is to think before sending sensitive material to a third party.

Suggestion

Try this fictional prompt: "Build a monthly budget with income of 3,000 units, housing of 1,000 and food of 400. Show the calculation and list missing categories." Replace details with rounded category totals only if needed.

Conclusion

Give AI enough to solve the task, while keeping the underlying financial record under your control.

Sources and further reading

[N1] Siladitya Ray. Forbes. 2 May 2023. Samsung Bans ChatGPT Among Employees After Sensitive Code Leak

02 Never buy certainty from a prediction

Personal / Finance and investing

A confident forecast can make a risky decision feel settled. AI-generated investment language is still a claim to investigate. A polished chart or an impressive success story does not establish how an investment will perform.

Do

Use AI to explain terms and prepare questions about fees, access to your money and possible losses. Verify the provider through the relevant regulator's register in your country. Read original disclosures and consider advice from a qualified professional.

Don't

Buy an investment because a bot ranks it first, a famous face appears to endorse it, or a promoter promises guaranteed returns through AI. Do not give a chatbot trading authority simply because its explanation sounds sophisticated.

Official guidance

A joint SEC, NASAA and FINRA investor alert warns about AI-branded investment fraud and unrealistic claims of guaranteed gains. [G2] The alert supports caution about the sales claim; it does not assess an investment you may be considering.

Suggestion

Ask: "Explain the general risks and fees of this investment type. Separate facts from assumptions, and list questions for a regulated adviser. Do not recommend a purchase." Check every factual answer against independent documents.

Conclusion

Use AI to improve your questions. Keep the decision tied to verified information, your circumstances and the possibility of loss.

Sources and further reading

[G2] SEC Office of Investor Education and Advocacy, NASAA and FINRA. 25 January 2024. Investor alert. Artificial Intelligence (AI) and Investment Fraud

03 Verify the person before moving the money

Personal / Money and identity

Urgency makes it harder to think. A call that sounds like a relative or a video that resembles your manager can still be false. A request for secrecy removes the people who could help you check it.

Do

Pause the transaction. Contact the person through a number or channel you already trust, obtained independently of the suspicious message. At work, follow the normal payment approval process even when the request appears to come from a senior person.

Don't

Send money, reveal a code or install remote-access software because a familiar voice says there is no time to check. Do not use the callback number supplied by the caller as your only verification.

Official guidance

The FTC explains that scammers can imitate a family member’s voice and recommends independently contacting the person. [G1] Familiarity is not reliable authentication.

Suggestion

Agree on a household response now: end the call, contact the person on a saved number and consult another trusted person. A private family phrase can add a check, but should not replace independent contact.

If you already paid

Contact your bank or payment provider immediately using its official app or published number. Explain the suspected fraud and ask about securing the account and possible recovery. Keep transaction details and report through your local fraud channel.

Conclusion

A genuine urgent request can withstand a brief, independent check.

Sources and further reading

[G1] Federal Trade Commission. March 2023. Consumer guidance. Scammers use AI to enhance their family emergency schemes

04 Let AI prepare questions for care

Personal / Health

An explanation can help you understand a medical term, but a chat response cannot provide the examination, full history and clinical judgment needed for your care. Even an answer with a reassuring tone may omit something important.

Do

Use general questions to prepare for an appointment. Ask a clinician or pharmacist to check advice about treatment, interactions and medication. Keep medical documents and other people's health details out of general chats unless an appropriate approved process protects them.

Don't

Start, stop or change medicine on a chatbot's instruction. Do not delay urgent care while trying different prompts. For an emergency, use local emergency services or urgent medical care.

News and guidance

AP reported that an expert found potentially dangerous omissions in an AI answer about wild mushrooms. [N3] WHO also warns that plausible AI health answers can contain serious errors. [G3] These examples support verification, not a diagnosis of any individual.

Suggestion

Try: "Help me make a short list of general questions to ask a pharmacist about a new prescription. Do not choose a dose or recommend changing treatment." Take the list to the professional and use their advice.

Conclusion

Use AI to arrive better prepared for human care, and keep medical decisions with qualified people who can assess you.

Sources and further reading

[N3] Associated Press. 31 May 2024. Google makes fixes to AI-generated search summaries after outlandish answers went viral

[G3] World Health Organization. 16 May 2023. Departmental update. WHO calls for safe and ethical AI for health

05 Do not outsource your support network

Personal / Wellbeing and relationships

A chatbot can help you find words for a difficult conversation. It may respond with warmth or agreement, but it cannot take the place of people who know you, notice changes and can help in the real world.

Do

Use a brief, non-identifying prompt to organize your thoughts. Make time for friends, family, a support group or a qualified mental health professional. Notice whether chatting is crowding out sleep, responsibilities or relationships.

Don't

Treat repeated agreement as proof that your interpretation is correct. Avoid letting a bot decide whether someone is trustworthy or whether you should end a relationship based on one description. Do not depend on it as your only crisis support.

Context and limits

UNICEF's family guidance discusses concerns about young people relying on AI instead of trusted people. [G4] That guidance is about children and teens; the adult boundary recommended here is a practical precaution, not a claim about a measured adult outcome.

Suggestion

Ask: "Help me draft a calm message asking someone I trust for a conversation. Do not judge either person or infer motives." Edit it so it sounds like you. Then actually contact the person.

When safety is urgent

If you may hurt yourself or cannot stay safe, contact local emergency services or a crisis service, and reach someone who can stay with you.

Conclusion

Let AI support a conversation that leads back to people.

Sources and further reading

[G4] UNICEF USA. 9 February 2026. Expert interview and family guidance. Protecting Children in the Age of AI

06 Read the policy behind the promise

Personal / Shopping and travel

A customer-service answer may sound official because it appears on a company website. Before you spend money or miss a deadline, check the actual policy and the terms that apply to your booking or purchase.

Do

Read the dated refund, cancellation or warranty terms. If an answer conflicts with them, get clarification from an authorized human representative before relying on it. Retain your receipt and the relevant correspondence.

Don't

Assume a chatbot can grant an exception, change a contract or guarantee a refund. Do not enter passport details or payment information into a chat unless it is part of a verified, appropriate transaction process.

News example

Ars Technica reported that Air Canada was ordered to provide a partial refund after its chatbot gave a passenger incorrect information about bereavement fares. [N4] This particular dispute does not establish what another company or court will do.

Suggestion

Use AI to list questions: “Which cancellation conditions should I confirm before buying a non-refundable ticket?” Then find the answers in the carrier’s current published terms and your specific fare conditions.

Conclusion

Let a chatbot point you toward the policy. Make costly decisions only after you have checked the policy itself.

Sources and further reading

[N4] Ashley Belanger. Ars Technica. 16 February 2024. Air Canada must honor refund policy invented by airline’s chatbot

07 Permission comes before the upload

Work / Confidential information

Being able to access a work document does not automatically mean you may send it to an outside service. Meeting notes can contain private information about colleagues, customers and future plans.

Do

Use the organization’s approved tool for the approved purpose and data type. Check who may access the result, how long data is retained, whether uploads are used for training, and whether connected services receive it. Ask the data owner if the rules are unclear.

Don't

Paste customer lists, staff records, confidential contracts, unreleased plans or proprietary code into a personal AI account. Do not assume that paying for a tool makes every upload acceptable.

News example

The Samsung incident described in Rule 1 shows why an employer may restrict uploads of internal material. [N1] It does not prove that all uploaded information becomes public. Controls differ by product, agreement and configuration.

Suggestion

For a meeting-summary template, use invented names and a fictional agenda. If real content is necessary, confirm the approved workflow first. Check the finished summary for missing decisions and incorrectly assigned actions.

If a mistake happens

Stop sharing and notify the relevant security or privacy contact promptly. Give the tool name, time and type of information shared. Follow their process; do not conceal the incident or assume deleting a chat resolves it.

Conclusion

Permission to use AI and permission to share a particular record are separate decisions.

Sources and further reading

[N1] Siladitya Ray. Forbes. 2 May 2023. Samsung Bans ChatGPT Among Employees After Sensitive Code Leak

08 Open every source that carries the argument

Work / Research and legal information

A citation can look complete while pointing to nothing, or to a real source that does not support the sentence beside it. Asking the same chatbot whether it is sure is not an independent check.

Do

Open the original source. Confirm its author, date, relevant passage and context. Check quotations word for word. For legal research, verify authorities in an authoritative legal source and have a qualified professional review consequential work.

Don't

Submit a report, legal document or public claim with references you have not checked. Do not substitute a plausible-looking link for evidence.

News example

AP reported sanctions against lawyers and their firm after court filings included nonexistent cases generated by ChatGPT. The judge also criticized how the lawyers responded when the citations were challenged. [N2] The lesson includes taking corrections seriously, not just checking the first draft.

Suggestion

Create a short claim log: the sentence, original source, supporting passage and verification date. Mark any unresolved point clearly. If you cannot verify a claim, remove it or state the uncertainty.

A useful prompt

“List the factual claims in this draft that need evidence. Do not invent references.” This creates a checking list; it does not complete the checking.

Conclusion

A claim is ready when its evidence has been inspected, not when its citation looks convincing.

Sources and further reading

[N2] Larry Neumeister. Associated Press, via News4JAX. 22 June 2023. Lawyers submitted bogus case law created by ChatGPT. A judge fined them $5,000

09 Keep a person accountable for every draft

Work / Writing and applications

A polished draft can hide a wrong date, a false promise or an invented achievement. The person approving the final work needs to understand what it says and be able to defend it.

Do

Use AI for outlines, alternative wording and editing. Check names, figures, dates, quotations and commitments against the underlying record. Follow your workplace, client or school rules on disclosure.

Don't

Send an unchecked output under your name. Do not add invented credentials to a CV, fabricated experience to a proposal or an unsupported allegation to a complaint. Avoid using AI to make your uncertainty look like expertise.

Related news

The fake-citation court case shows that responsibility remained with the people who submitted the material. [N2] Applying that lesson to routine writing is this book's recommendation; it is not a claim that the same legal outcome applies everywhere.

Suggestion

Review a draft in two passes. First, check every fact and promise. Second, check whether the tone, level of certainty and intended audience are right. Keep a brief record of substantial AI assistance where your rules require it.

Fictional example

Replace "I led a team of ten" with your true role, even if the invented version reads better. A stronger sentence cannot repair a false statement.

Conclusion

Use AI to improve how you express your work while keeping the content accurate and yours to own.

Sources and further reading

[N2] Larry Neumeister. Associated Press, via News4JAX. 22 June 2023. Lawyers submitted bogus case law created by ChatGPT. A judge fined them $5,000

10 Give automation only the access it needs

Work / Tools and automation

An AI that drafts text has different consequences from an AI that can send messages, delete files or spend money. Decide what it may do before you connect an account.

Do

Start with the smallest task and the fewest permissions. Prefer previews and human confirmation for consequential actions. Test on copies or invented data, keep backups, and understand how to revoke access.

Don't

Give broad mailbox, storage, payment or administrator access simply for convenience. Do not let instructions inside an email, webpage or document authorize new actions. Material the AI reads may come from someone you do not trust.

Practical context

This is a precautionary workflow recommendation, not a summary of a particular news incident. The important question is concrete: if this tool misunderstands the request, what can it change before you notice?

Suggestion

For an email assistant, begin with drafting only. Before allowing sending, verify the recipient, attachments and final text each time. For a file task, use a copied folder and compare the changes before applying them to originals.

Recovery plan

Write down where to disconnect the tool, who can restore affected data and who must be told if it acts incorrectly. Check the actual service’s current help pages rather than trusting a generated menu path.

Conclusion

Expand access only when the benefit is clear and you can review or recover the result.

Sources and further reading

This rule offers original practical guidance; no specific news incident is cited.

11 Check the claim before you amplify it

Public life / News and social media

Images, voices and fluent summaries can all make a weak claim feel real. A post may also be old material presented as new. Repeating a claim to criticize it can spread it further.

Do

Find the original publisher and date. Look for independent reporting or a relevant primary record. Verify that the source actually supports the claim. If you cannot establish what happened, avoid presenting it as fact.

Don't

Forward dramatic screenshots or AI summaries on appearance alone. Do not treat the absence of visual flaws as proof that a video is authentic. Avoid relying on a single AI detector as a final verdict.

News example

AP reported both erroneous AI search answers and fabricated screenshots that falsely purported to show AI answers. [N3] Verification matters on both sides: a chatbot's claim and a viral claim about a chatbot.

Suggestion

Before sharing, answer three questions in your own words: Who first published this? What evidence can I inspect? What remains uncertain? If one answer is missing, wait or add a clear qualification instead of guessing.

If you shared a mistake

Correct it in the same place if possible. Say what was wrong and link to reliable information. Tell people directly if the error could affect a decision.

Conclusion

The speed of sharing should follow the strength of the evidence.

Sources and further reading

[N3] Associated Press. 31 May 2024. Google makes fixes to AI-generated search summaries after outlandish answers went viral

12 Ask before borrowing a face or voice

Public life / Images and consent

Keep it private

Do not casually upload photos, videos, screenshots, messages or files you want to keep private. For an AI experiment, use a fictional example or a non-sensitive description. An upload is still sharing with a service even if you never publish the result.

An AI edit can make someone appear to say or do something they never agreed to. The harm can continue after the creator calls it a joke. A label does not replace the person's permission.

Do

Obtain clear permission for the specific use of a person's image or voice. Explain the intended audience and whether it will be published. Use fictional people or appropriately licensed material when permission is absent.

Don't

Create deceptive endorsements, impersonations or humiliating images of real people. Never create or share sexual images involving minors. Do not upload somebody else's private photo collection or messages to experiment with a tool.

Official context

Australia's eSafety Commissioner describes the impact of AI-generated image abuse on school communities. [G5] Its reporting services and legal details are jurisdiction-specific; the consent principle here is a broader ethical recommendation.

Suggestion

For a fictional advertisement, use an invented character and clearly identify synthetic media when it could otherwise mislead. Before commercial use, check the tool's current terms and relevant rights, and seek legal advice when necessary.

If someone is targeted

Support them, report through the platform and seek appropriate local help. Avoid forwarding harmful images as evidence; record links, account names and dates instead, especially when minors are involved.

Conclusion

Creative freedom works best when the people depicted retain dignity and meaningful choice.

Sources and further reading

[G5] Australian eSafety Commissioner. Safety guidance. Accessed 7 October 2026. Deepfake damage in schools

13 Agree on boundaries before there is a problem

Family / Supporting teenagers

Children's photos

Avoid unnecessary uploads or forwarding of children's pictures, including to AI filters, avatars and novelty edits. Use an imaginary character instead. If sharing is necessary, check the purpose, audience and service settings; obtain appropriate parent or guardian permission and involve the child in the decision. A private chat can still be copied or reshared. [G6]

Young people need clear rules and an adult they can approach without fear of immediate blame. A household plan works better when it explains both the benefits of AI and the situations that require help.

Do

Review age requirements, privacy settings and school rules together. Agree on appropriate uses, spending limits and when to stop. Name more than one trusted adult a teen can approach. Revisit the plan as the tools and the teen’s needs change.

Don't

Treat parental controls as a guarantee of safety. Do not tell a young person to rely on a bot for emergencies or keep worrying conversations secret. Avoid responding to a disclosure with blame that makes them less likely to ask next time.

Official context

UNICEF’s family guidance encourages discussing AI, privacy and balanced use with children and teens. [G4] A useful starting point is an open conversation about what the young person enjoys and what feels uncomfortable.

Suggestion

Create a short agreement: “We protect private information, check school rules, stop when something feels wrong, and ask a person for help.” Include a regular check-in and make clear that the young person can ask for help after a mistake.

Conclusion

The most useful safety rule is one a teenager understands and feels able to follow in a difficult moment.

Sources and further reading

[G6] Australian eSafety Commissioner. Parent guidance. Accessed 13 October 2026. Privacy and your child

[G4] UNICEF USA. 9 February 2026. Expert interview and family guidance. Protecting Children in the Age of AI

14 Measure the result before trusting the shortcut

Everyday practice / Choosing useful tools

An AI workflow only saves time if its output is good enough after checking. A fast answer that requires extensive repair may be a poor fit for the job.

Do

Define success before you begin: accurate figures, accessible language, a usable plan or time saved after review. Try a small, low-risk task. Compare the result with an existing trusted method and record what had to be corrected.

Don't

Connect more accounts or buy more access just because a demonstration looks impressive. Do not move a workflow into health, hiring, credit or other consequential decisions without appropriate expertise and safeguards.

Practical context

This rule is an evaluation method, not a claim from a particular news story. Useful questions include whether the tool fails predictably, whether you can detect its errors and whether a simpler method would meet your needs.

Suggestion

For a weekly planning task, compare three attempts with your normal approach. Count review time as well as drafting time. If the system repeatedly invents commitments, narrow its task to formatting a list you already checked.

Keep the skill

Regularly perform part of the task yourself. You need enough understanding to spot a mistake, explain the result and continue working if the service is unavailable.

Conclusion

Adopt AI where you can demonstrate a benefit and maintain the judgment needed to supervise it.

Sources and further reading

This rule offers original practical guidance; no specific news incident is cited.

Your check before sending

Before the prompt	What am I trying to achieve? Can I use invented details or a shorter summary? Do I have permission to share every part of the material, including information about other people?
Before trusting the answer	Which statements are factual? Have I opened the original sources? Do the dates, calculations and quotations check out? Have I separated uncertainty from evidence?
Before an action	Could this affect money, health, legal rights, employment, reputation or safety? Who needs to review it? Can I undo the action? Is the tool asking for more access than the task requires?
When something goes wrong	Stop the activity. Preserve relevant non-sensitive details, such as times, links and transaction records. Contact the appropriate person: a bank, clinician, manager, security team or platform support. Change exposed credentials through the official service. Do not assume deleting one chat removes every copy.
A repeatable prompt	“Help me with [specific task] using these non-sensitive details. Ask about missing information rather than inventing it. Separate facts, assumptions and suggestions. Flag claims I should verify and actions that need human review.”
Remember	A careful prompt reduces ambiguity. It does not guarantee a safe or accurate response. Your final check still matters.

Conclusion

You do not need to reject AI or trust it completely. You can choose a useful task, set limits and judge the output against evidence. That habit matters more than any single prompt.

Start with one change

Choose the rule most relevant to your next task. If you use AI for money questions, replace full documents with non-identifying summaries. If you use it for research, open every important reference. If you connect it to other tools, reduce permissions and require review for consequential actions.

Keep three boundaries

Private data needs a justified purpose and appropriate protection. Important claims need checked evidence. Decisions with serious consequences need accountable human judgment.

Review your practice

Once a month, revisit which tools you use, what they can access and whether they are still helping. Remove connections you no longer need. Check current product settings rather than assuming last year’s policy still applies.

Your next step

Write down one task AI can help with and one boundary you will apply. Carry out the task, check the result and decide whether the tool earned a place in your routine. The goal is useful assistance that you understand and can control.